

## Importing the Dependencies

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

## Data Collection and Processing

```
# loading the csv data to a Pandas DataFrame
heart_data = pd.read_csv('/content/data.csv')
```

```
# print first 5 rows of the dataset
heart_data.head()
```

```
➡
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1

```
# print last 5 rows of the dataset
heart_data.tail()
```

```
➡
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
298	57	0	0	140	241	0	1	123	1	0.2	1	0	3	0
299	45	1	3	110	264	0	1	132	0	1.2	1	0	3	0
300	68	1	0	144	193	1	1	141	0	3.4	1	2	3	0
301	57	1	0	130	131	0	1	115	1	1.2	1	1	3	0
302	57	0	1	130	236	0	0	174	0	0.0	1	1	2	0

```
# number of rows and columns in the dataset
heart_data.shape
```

```
➡ (303, 14)
```

```
# getting some info about the data
heart_data.info()
```

```
➡ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
---
```

```

0   age      303 non-null   int64
1   sex      303 non-null   int64
2   cp       303 non-null   int64
3   trestbps 303 non-null   int64
4   chol     303 non-null   int64
5   fbs      303 non-null   int64
6   restecg  303 non-null   int64
7   thalach  303 non-null   int64
8   exang    303 non-null   int64
9   oldpeak  303 non-null   float64
10  slope    303 non-null   int64
11  ca       303 non-null   int64
12  thal     303 non-null   int64
13  target   303 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB

```

```

# checking for missing values
heart_data.isnull().sum()

```

```

⇒ age      0
sex      0
cp       0
trestbps 0
chol     0
fbs      0
restecg  0
thalach  0
exang    0
oldpeak  0
slope    0
ca       0
thal     0
target   0
dtype: int64

```

```

# statistical measures about the data
heart_data.describe()

```

```

⇒
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	
<b>count</b>	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.
<b>mean</b>	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515	0.528053	149.646865	0.
<b>std</b>	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198	0.525860	22.905161	0.
<b>min</b>	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.000000	71.000000	0.
<b>25%</b>	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000	0.000000	133.500000	0.
<b>50%</b>	55.000000	1.000000	1.000000	130.000000	240.000000	0.000000	1.000000	153.000000	0.
<b>75%</b>	61.000000	1.000000	2.000000	140.000000	274.500000	0.000000	1.000000	166.000000	1.
<b>max</b>	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	2.000000	202.000000	1.

```

# checking the distribution of Target Variable
heart_data['target'].value_counts()

```

```

⇒ 1    165
   0    138
   Name: target, dtype: int64

```

1 --> Defective Heart

0 --> Healthy Heart

### Splitting the Features and Target

```

X = heart_data.drop(columns='target', axis=1)
Y = heart_data['target']

```

```
print(X)
```

```

⇒   age  sex  cp  trestbps  chol  ...  exang  oldpeak  slope  ca  thal
0    63   1   3      145   233  ...    0      2.3      0  0    1
1    37   1   2      130   250  ...    0      3.5      0  0    2
2    41   0   1      130   204  ...    0      1.4      2  0    2
3    56   1   1      120   236  ...    0      0.8      2  0    2
4    57   0   0      120   354  ...    1      0.6      2  0    2
..  ...  ...  ..      ...   ...  ...  ...    ...    ...    ...  ..  ...
298  57   0   0      140   241  ...    1      0.2      1  0    3
299  45   1   3      110   264  ...    0      1.2      1  0    3
300  68   1   0      144   193  ...    0      3.4      1  2    3
301  57   1   0      130   131  ...    1      1.2      1  1    3
302  57   0   1      130   236  ...    0      0.0      1  1    2

```

[303 rows x 13 columns]

```
print(Y)
```

```

⇒ 0    1
   1    1
   2    1
   3    1
   4    1
   ..
298  0
299  0
300  0
301  0
302  0
   Name: target, Length: 303, dtype: int64

```

### Splitting the Data into Training data & Test Data

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, stratify=Y, random_state=
```

```
print(X.shape, X_train.shape, X_test.shape)
```

```

⇒ (303, 13) (242, 13) (61, 13)

```

### Model Training

## Logistic Regression

```
model = LogisticRegression()
```

```
# training the LogisticRegression model with Training data
model.fit(X_train, Y_train)
```

```
➦ /usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py:940: ConvergenceWarning:
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
                    warm_start=False)
```

## Model Evaluation

### Accuracy Score

```
# accuracy on training data
X_train_prediction = model.predict(X_train)
training_data_accuracy = accuracy_score(X_train_prediction, Y_train)
```

```
print('Accuracy on Training data : ', training_data_accuracy)
```

```
➦ Accuracy on Training data : 0.8512396694214877
```

```
# accuracy on test data
X_test_prediction = model.predict(X_test)
test_data_accuracy = accuracy_score(X_test_prediction, Y_test)
```

```
print('Accuracy on Test data : ', test_data_accuracy)
```

```
➦ Accuracy on Test data : 0.819672131147541
```

## Building a Predictive System


```
input_data = (62,0,0,140,268,0,0,160,0,3.6,0,2,2)
```

```
# change the input data to a numpy array
input_data_as_numpy_array= np.asarray(input_data)
```

```
# reshape the numpy array as we are predicting for only on instance
input_data_reshaped = input_data_as_numpy_array.reshape(1,-1)
```

```
prediction = model.predict(input_data_reshaped)
print(prediction)
```

```
if (prediction[0]== 0):  
    print('The Person does not have a Heart Disease')  
else:  
    print('The Person has Heart Disease')
```

 [0]  
The Person does not have a Heart Disease